A Course in Reinforcement Learning

2nd Edition

Dimitri P. Bertsekas

Athena Scientific

A Course in Reinforcement Learning (2nd Edition)

by

Dimitri P. Bertsekas

Arizona State University

WWW site for book information and orders

http://www.athenasc.com

Athena Scientific, Belmont, Massachusetts

Athena Scientific
Post Office Box 805
Nashua, NH 03060
U.S.A.

Email: info@athenasc.com
WWW: http://www.athenasc.com

Cover design by Dimitri P. Bertsekas

© 2026 Dimitri P. Bertsekas
All rights reserved. No part of this book may be reproduced in any form by any electronic or mechanical means (including photocopying, recording, or information storage and retrieval) without permission in writing from the publisher.

Publisher's Cataloging-in-Publication Data

Bertsekas, Dimitri P.
A Course in Reinforcement Learning (2nd Edition)
Includes Bibliography and Index
1. Mathematical Optimization. 2. Dynamic Programming. I. Title.
QA402.5 .B465 2020 519.703 00-91281

ISBN-10: 1-886529-29-9, ISBN-13: 978-1-886529-29-8

Latest editorial and other changes added on 12/26/2025

ABOUT THE AUTHOR

Dimitri Bertsekas studied Mechanical and Electrical Engineering at the National Technical University of Athens, Greece, and obtained his Ph.D. in system science from the Massachusetts Institute of Technology. He has held faculty positions with the Engineering-Economic Systems Department, Stanford University, and the Electrical Engineering Department of the University of Illinois, Urbana. From 1979 to 2019 he was a professor at the Electrical Engineering and Computer Science Department of the Massachusetts Institute of Technology (M.I.T.), where he continues to hold the title of McAfee Professor of Engineering. In 2019, he joined the School of Computing and Augmented Intelligence at the Arizona State University, Tempe, AZ, as Fulton Professor of Computational Decision Making.

Professor Bertsekas' teaching and research have spanned several fields, including deterministic optimization, dynamic programming and stochastic control, large-scale and distributed computation, artificial intelligence, and data communication networks. He has authored or coauthored numerous research papers and twenty books, several of which are currently used as textbooks in MIT classes, including "Dynamic Programming and Optimal Control," "Data Networks," "Introduction to Probability," and "Nonlinear Programming." At ASU, he has been focusing in teaching and research in reinforcement learning, and he has written several textbooks and research monographs in this field since 2019.

Professor Bertsekas was awarded the INFORMS 1997 Prize for Research Excellence in the Interface Between Operations Research and Computer Science for his book "Neuro-Dynamic Programming" (co-authored with John Tsitsiklis), the 2001 AACC John R. Ragazzini Education Award, the 2009 INFORMS Expository Writing Award, the 2014 AACC Richard Bellman Heritage Award, the 2014 INFORMS Khachiyan Prize for Life-Time Accomplishments in Optimization, the 2015 MOS/SIAM George B. Dantzig Prize, and the 2022 IEEE Control Systems Award. In 2018 he shared with his coauthor, John Tsitsiklis, the 2018 INFORMS John von Neumann Theory Prize for the contributions of the research monographs "Parallel and Distributed Computation" and "Neuro-Dynamic Programming." Professor Bertsekas was elected in 2001 to the United States National Academy of Engineering for "pioneering contributions to fundamental research, practice and education of optimization/control theory."

ATHENA SCIENTIFIC

OPTIMIZATION AND COMPUTATION SERIES

1. A Course in Reinforcement Learning, 2nd Edition, by Dimitri P. Bertsekas, 2025, ISBN 978-1-886529-29-8, 510 pages
2. Lessons from AlphaZero for Optimal, Model Predictive, and Adaptive Control by Dimitri P. Bertsekas, 2022, ISBN 978-1-886529-17-5, 245 pages
3. Abstract Dynamic Programming, 3rd Edition, by Dimitri P. Bertsekas, 2022, ISBN 978-1-886529-47-2, 420 pages
4. Rollout, Policy Iteration, and Distributed Reinforcement Learning, by Dimitri P. Bertsekas, 2020, ISBN 978-1-886529-07-6, 480 pages
5. Reinforcement Learning and Optimal Control, by Dimitri P. Bertsekas, 2019, ISBN 978-1-886529-39-7, 388 pages
6. Dynamic Programming and Optimal Control, Two-Volume Set, by Dimitri P. Bertsekas, 2017, ISBN 1-886529-08-6, 1270 pages
7. Nonlinear Programming, 3rd Edition, by Dimitri P. Bertsekas, 2016, ISBN 1-886529-05-1, 880 pages
8. Convex Optimization Algorithms, by Dimitri P. Bertsekas, 2015, ISBN 978-1-886529-28-1, 576 pages
9. Convex Optimization Theory, by Dimitri P. Bertsekas, 2009, ISBN 978-1-886529-31-1, 256 pages
10. Introduction to Probability, 2nd Edition, by Dimitri P. Bertsekas and John N. Tsitsiklis, 2008, ISBN 978-1-886529-23-6, 544 pages
11. Convex Analysis and Optimization, by Dimitri P. Bertsekas, Angelia Nedić, and Asuman E. Ozdaglar, 2003, ISBN 1-886529-45-0, 560 pages
12. Network Optimization: Continuous and Discrete Models, by Dimitri P. Bertsekas, 1998, ISBN 1-886529-02-7, 608 pages
13. Network Flows and Monotropic Optimization, by R. Tyrrell Rockafellar, 1998, ISBN 1-886529-06-X, 634 pages
14. Introduction to Linear Optimization, by Dimitris Bertsimas and John N. Tsitsiklis, 1997, ISBN 1-886529-19-1, 608 pages
15. Parallel and Distributed Computation: Numerical Methods, by Dimitri P. Bertsekas and John N. Tsitsiklis, 1997, ISBN 1-886529-01-9, 718 pages
16. Neuro-Dynamic Programming, by Dimitri P. Bertsekas and John N. Tsitsiklis, 1996, ISBN 1-886529-10-8, 512 pages
17. Constrained Optimization and Lagrange Multiplier Methods, by Dimitri P. Bertsekas, 1996, ISBN 1-886529-04-3, 410 pages
18. Stochastic Optimal Control: The Discrete-Time Case, by Dimitri P. Bertsekas and Steven E. Shreve, 1996, ISBN 1-886529-03-5, 330 pages

CONTENTS

1. Exact and Approximate Dynamic Programming

1.1. AlphaZero, Off-Line Training, and On-Line Play . . . . . p. 4
1.2. Deterministic Dynamic Programming . . . . . p. 10
1.2.1. Finite Horizon Problem Formulation . . . . . p. 10
1.2.2. The Dynamic Programming Algorithm . . . . . p. 14
1.2.3. Approximation in Value Space and Rollout . . . . . p. 22
1.3. Stochastic Exact and Approximate Dynamic Programming . . p. 28
1.3.1. Finite Horizon Problems . . . . . p. 28
1.3.2. Approximation in Value Space for Stochastic DP . . . p. 34
1.3.3. Approximation in Policy Space . . . . . p. 38
1.3.4. Training of Cost Function and Policy . . . . . Approximations . . . . . p. 41
1.4. Infinite Horizon Problems - An Overview . . . . . p. 42
1.4.1. Infinite Horizon Methodology . . . . . p. 46
1.4.2. Approximation in Value Space - Infinite Horizon . . . p. 49
1.4.3. Understanding Approximation in Value Space . . . . . p. 55
1.5. Newton's Method - Linear Quadratic Problems . . . . . p. 56
1.5.1. Visualizing Approximation in Value Space - Region of . . . Stability . . . . . p. 62
1.5.2. Rollout and Policy Iteration . . . . . p. 70
1.5.3. Local and Global Error Bounds for Approximation in . . . Value Space . . . . . p. 74
1.5.4. How Approximation in Value Space Can Fail and . . . . . What to Do About It . . . . . p. 78
1.6. Examples, Reformulations, and Simplifications . . . . . p. 79
1.6.1. A Few Words About Modeling . . . . . p. 79
1.6.2. Problems with a Termination State . . . . . p. 83
1.6.3. General Discrete Optimization Problems . . . . . p. 85
1.6.4. General Finite to Infinite Horizon Reformulation . . . p. 89
1.6.5. State Augmentation, Time Delays, Forecasts, and . . . . . Uncontrollable State Components . . . . . p. 91
1.6.6. Partial State Information and Belief States . . . . . p. 97
1.6.7. Multiagent Problems and Multiagent Rollout . . . . p. 102
1.6.8. Problems with Unknown or Changing Model - . . . . . Adaptive Control . . . . . p. 107
1.6.9. Model Predictive Control . . . . . p. 117
1.7. Reinforcement Learning and Decision/Control . . . . . p. 131
1.7.1. Differences in Terminology . . . . . p. 131
1.7.2. Differences in Notation . . . . . p. 133
1.7.3. Relations Between DP and RL . . . . . p. 134

1.7.4. Synergy Between Large Language Models and RL . p. 135
1.7.5. Machine Learning and Optimization . . . . . p. 136
1.7.6. Mathematics in Machine Learning and Optimization p. 139
1.8. Notes, Sources, and Exercises . . . . . p. 140

2. Approximation in Value Space - Rollout Algorithms

2.1. Deterministic Finite Horizon Problems . . . . . p. 166
2.2. Approximation in Value Space - Deterministic Problems . . p. 173
2.3. Rollout Algorithms for Discrete Optimization . . . . . p. 177
2.3.1. Cost Improvement with Rollout - Sequential Consistency, . . Sequential Improvement . . . . . p. 183
2.3.2. The Fortified Rollout Algorithm . . . . . p. 189
2.3.3. Using Multiple Base Heuristics - Parallel Rollout . . . p. 192
2.3.4. Using Multiple Rollout Trajectories - General Rollout . p. 194
2.3.5. Simplified Rollout Algorithms . . . . . p. 195
2.3.6. Truncated Rollout with Terminal Cost Approximation . p. 196
2.3.7. Rollout with an Expert - Model-Free Rollout . . . . . p. 197
2.3.8. Using a World Model for Rollout . . . . . p. 202
2.3.9. Local Search with Rollout for Discrete Optimization . . p. 204
2.3.10. Inference in n-Grams, Transformers, HMMs, and . . . . . Markov Chains . . . . . p. 208
2.4. Approximation in Value Space with Multistep Lookahead . p. 221
2.4.1. Iterative Deepening Using Forward Dynamic . . . . . Programming . . . . . p. 227
2.4.2. Incremental Multistep Rollout . . . . . p. 229
2.5. Constrained Forms of Rollout Algorithms . . . . . p. 232
2.5.1. Constrained Rollout for Discrete Optimization and Integer . . Programming . . . . . p. 245
2.6. Small Stage Costs and Long Horizon - Continuous-Time . . . . . Rollout . . . . . p. 249
2.7. Approximation in Value Space - Stochastic Problems . . . p. 257
2.7.1. Simplified Rollout and Policy Iteration . . . . . p. 262
2.7.2. Certainty Equivalence Approximations . . . . . p. 261
2.7.3. Simulation-Based Implementation of the Rollout . . . . . Algorithm . . . . . p. 265
2.7.4. Variance Reduction in Rollout - Comparing Advantages p. 267
2.7.5. Monte Carlo Tree Search . . . . . p. 270
2.7.6. Randomized Policy Improvement by Monte Carlo . . . . . Tree Search . . . . . p. 274
2.8. Rollout for Infinite-Spaces Problems - Optimization . . . . . Heuristics . . . . . p. 274
2.8.1. Rollout for Infinite-Spaces Deterministic Problems . . . p. 275
2.8.2. Rollout Based on Stochastic Programming . . . . . p. 279
2.8.3. Stochastic Rollout with Certainty Equivalence . . . . . p. 281

2.9. Multiagent Rollout . . . . . p. 283
2.9.1. Asynchronous and Autonomous Multiagent Rollout . . p. 294
2.10. Rollout for Bayesian Optimization and Sequential . . . . . Estimation . . . . . p. 297
2.11. Adaptive Control by Rollout with a POMDP Formulation . p. 309
2.12. Minimax Control and Reinforcement Learning . . . . . p. 317
2.12.1. Exact Dynamic Programming for Minimax Problems . p. 318
2.12.2. Minimax Approximation in Value Space and Rollout . p. 321
2.12.3. Combined Approximation in Value and Policy Space . . for Minimax Control . . . . . p. 327
2.12.4. A Meta Algorithm for Computer Chess Based on . . Reinforcement Learning . . . . . p. 328
2.12.5. Combined Approximation in Value and Policy Space . . for Sequential Noncooperative Games . . . . . p. 332
2.13. Notes, Sources, and Exercises . . . . . p. 334

3. Learning Values and Policies . . . . . .

3.1. Parametric Approximation Architectures . . . . . p. 349
3.1.1. Cost Function Approximation . . . . . p. 350
3.1.2. Feature-Based Architectures . . . . . p. 351
3.1.3. Training of Linear and Nonlinear Architectures . . . p. 362
3.2. Neural Networks . . . . . p. 370
3.2.1. Training of Neural Networks . . . . . p. 375
3.2.2. Multilayer and Deep Neural Networks . . . . . p. 376
3.3. Learning Cost Functions in Approximate DP . . . . . p. 377
3.3.1. Fitted Value Iteration . . . . . p. 378
3.3.2. Q-Factor Parametric Approximation - Model-Free . . . . Implementation . . . . . p. 380
3.3.3. Parametric Approximation in Infinite Horizon . . . . . Problems - Approximate Policy Iteration . . . . . p. 382
3.3.4. Optimistic Policy Iteration with Parametric Q-Factor . . . Approximation - SARSA and DQN . . . . . p. 386
3.3.5. Approximate Policy Iteration for Infinite Horizon . . . . . POMDP . . . . . p. 388
3.3.6. Advantage Updating - Approximating Q-Factor . . . . . Differences . . . . . p. 393
3.3.7. Differential Training of Cost Differences for Rollout . p. 395
3.4. Learning a Policy in Approximate DP . . . . . p. 397
3.4.1. The Use of Classifiers for Approximation in Policy . . . . Space . . . . . p. 398
3.4.2. Policy Iteration with Value and Policy Networks . . p. 402
3.4.3. Why Use On-Line Play and not Just Train a Policy . . . . Network to Emulate the Lookahead Minimization? . p. 404
3.5. Policy Gradient and Related Methods . . . . . p. 406

3.5.1. Gradient Methods for Parametric Cost Optimization . . . p. 407
3.5.2. Policy Gradient-Like Methods . . . p. 410
3.5.3. Scaling and Proximal Policy Optimization Methods . . . p. 415
3.5.4. Random Direction Methods . . . p. 422
3.5.5. Random Search and Cross-Entropy Methods . . . p. 424
3.5.6. Refining and Retraining Parametric Policies . . . p. 426
3.6. Aggregation . . . p. 427
3.6.1. Aggregation with Representative States . . . p. 428
3.6.2. Continuous Control Space Discretization . . . p. 435
3.6.3. Continuous State Space - POMDP Discretization . . . p. 436
3.6.4. General Aggregation . . . p. 438
3.6.5. Types of Aggregation and Error Bounds . . . p. 441
3.6.6. Aggregation Using Features . . . p. 443
3.6.7. Biased Aggregation . . . p. 444
3.6.8. Asynchronous Distributed Aggregation . . . p. 447
3.7. Notes, Sources, and Exercises . . . p. 449

References . . . p. 455