

Every chapter so far assumes access to a scalar reward signal: dynamic programming requires $r(s, a)$, model-free RL observes r_t after each transition, and the Bellman optimality equation presupposes that rewards are known or observable. When they are neither, the DP/RL machinery cannot be applied directly.

In many domains, scalar rewards are unavailable but ordinal preferences over trajectories are easy to elicit. A human evaluator cannot assign a meaningful numerical score to a paragraph of text, but can reliably say “response A is better than response B.” The raw data is a set of trajectory pairs with binary preference labels. Reinforcement learning from human feedback (RLHF) uses these ordinal comparisons to learn a proxy reward function r_θ , which then serves as the scalar signal for standard RL optimization. This is not an inverse problem in the IRL sense; the goal is not to rationalize observed behavior, but rather a two-stage forward problem in which the analyst first learns a reward from human judgments and then solves the resulting MDP. Christiano et al. (2017) demonstrated that this approach could train agents without an explicit reward function. RLHF has since become the predominant method for aligning large language models.

0.1 Learning Rewards from Preferences

The canonical RLHF framework is built on a formal model of human preference. The observed data consists of tuples (s, y_w, y_l) , where s is a context, and y_w and y_l are two outputs, with y_w being the “winner” preferred by a human. Assuming preferences follow a latent utility model, the Bradley-Terry model (Bradley and Terry, 1952) gives the probability that y_w is preferred: $P(y_w \succ y_l | s) = \sigma(r_\theta(s, y_w) - r_\theta(s, y_l))$, where $r_\theta : \mathcal{S} \times \mathcal{Y} \rightarrow \mathbb{R}$ is a learned *reward model* parameterized by θ , trained to approximate human preferences (not the ground-truth reward, which is unobserved), and $\sigma(\cdot)$ is the logistic function. This formulation is a binary logit model (Section ??, Equation ??).¹ The reward model parameters θ are estimated by minimizing the negative log-likelihood of the observed human choices:

$$\mathcal{L}(\theta) = -\mathbb{E}_{(s, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r_\theta(s, y_w) - r_\theta(s, y_l))], \quad (1)$$

¹Iskhakov et al. (2020) discuss the contrasts and synergies between machine learning and structural econometrics, including the shared reliance on logit-based choice models that underlies both RLHF and dynamic discrete choice estimation.

In the LLM setting, the “outputs” y_w and y_l are token sequences, i.e., trajectories of the autoregressive policy², so preferences are over trajectories rather than single actions. The preference loss in Equation (1) is MLE of a choice model where the “alternatives” are trajectory segments and the “choice” is the human-preferred one.

This logit foundation is also a social-choice assumption once labelers are heterogeneous. Pooling pairwise comparisons from many evaluators does not recover a neutral “human preference”; it estimates a single score whose odds ratios summarize the aggregation rule induced by the sampling scheme, label model, and loss function. Recent alignment work makes this point explicit: standard Bradley-Terry-style reward learning can behave like a hidden social welfare function, can fail natural axioms for aggregating rankings, and can create strategic reporting incentives when preferences differ across contexts or groups (Conitzer et al., 2024; Siththaranjan et al., 2024; Ge et al., 2024; Park et al., 2024). DPO inherits the same aggregation problem because it replaces the explicit reward model with an implicit reward difference, not with a theory of whose preferences count.

0.2 The RLHF Pipeline and Direct Optimization

The learned r_θ then serves as a proxy objective for policy optimization. Building on the reward-learning and RL fine-tuning framework of Ziegler et al. (2019) and Stiennon et al. (2020), the canonical three-stage pipeline was formalized by Ouyang et al. (2022). First, a base pretrained model is initialized via supervised fine-tuning (SFT)³ on a small dataset of high-quality demonstrations, yielding an initial policy π^{SFT} . This grounds the model in the desired style and format. The second step is the reward model training as described, using preference data generated from this π^{SFT} .

In the third step, the SFT policy π_ϕ is fine-tuned via PPO (Section ??) to maximize the frozen reward model r_θ ,⁴ with a KL-divergence penalty preventing the policy from

²An *autoregressive* language model generates text token by token: at each step it outputs a distribution over the vocabulary conditioned on all preceding tokens, then samples the next token; the full response is therefore a trajectory in token space and the model acts as a sequential policy over a vocabulary-sized action set.

³*Pretraining* optimizes a language model to predict the next token across a massive text corpus, producing broad linguistic knowledge with no behavioral objective. *Supervised fine-tuning* (SFT) continues training on a small curated dataset of (prompt, ideal-response) pairs to specialize the model toward the desired task and establish the reference policy π^{SFT} from which the KL penalty is measured.

⁴After fine-tuning concludes, the reward model may be reused for evaluation or for later alignment iterations, but it is not updated inside the PPO optimization loop.

drifting into regions where r_θ is unreliable. The objective is

$$J(\phi) = \mathbb{E}_{s \sim \mathcal{D}, y \sim \pi_\phi(\cdot|s)}[r_\theta(s, y)] - \lambda_{KL} \mathbb{E}_{s \sim \mathcal{D}}[D_{KL}(\pi_\phi(\cdot|s) || \pi^{SFT}(\cdot|s))], \quad (2)$$

where D_{KL} is the Kullback-Leibler divergence and λ_{KL} controls the penalty strength.⁵ Without this constraint, the policy exploits inaccuracies in r_θ to achieve high proxy scores with degenerate behavior (“reward hacking”), the RLHF analogue of divergence under function approximation (Section ??).

The KL-regularized objective in Equation (2) admits a Bayesian interpretation (Korbak et al., 2022). In this view, the reference policy π^{SFT} acts as a prior distribution over plausible responses. The reward model r_θ provides evidence, specifying which responses are more desirable. The goal of alignment is to find the posterior distribution π^* that optimally combines the prior with this evidence. This ideal posterior policy takes the form of Equation (3):

$$\pi^*(y|s) \propto \pi^{SFT}(y|s) \exp\left(\frac{r_\theta(s, y)}{\lambda_{KL}}\right). \quad (3)$$

The reward function scaled by λ_{KL} defines the log-likelihood, so the KL-regularized objective $J(\phi)$ is equivalent (up to an additive constant) to the Evidence Lower Bound (ELBO) for this Bayesian inference problem. Maximizing the RLHF objective via PPO is therefore variational inference: finding the policy π_ϕ that minimizes KL divergence to π^* . This reframes the KL penalty as a structural component of the inference problem rather than an ad-hoc regularizer.

Despite this closed-form characterization, the three-stage RLHF pipeline is complex to implement, requiring training multiple large models and a computationally expensive RL loop. Direct Preference Optimization (DPO), introduced by Rafailov et al. (2023), collapses the pipeline into a single supervised learning objective by reparameterizing the reward function in terms of π^* and π^{SFT} , as in Equation (4):

$$r(s, y) = \lambda_{KL} \log\left(\frac{\pi^*(y|s)}{\pi^{SFT}(y|s)}\right) + \lambda_{KL} \log Z(s). \quad (4)$$

When this analytical expression for the reward is substituted into the Bradley-Terry preference loss from Equation (1), the unknown partition function $Z(s)$ cancels out. This

⁵ λ_{KL} denotes the KL penalty weight, reserving β for model parameters and γ for discount factors. The standard RLHF literature, including Rafailov et al. (2023), uses β for this parameter.

yields a loss function that depends only on the policy π_ϕ being optimized and the fixed reference policy π^{SFT} . The DPO objective, given in Equation (5), is thus a simple binary cross-entropy loss over policy likelihoods:

$$\mathcal{L}_{DPO}(\phi; \pi^{SFT}) = -\mathbb{E}_{(s, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\lambda_{KL} \log \frac{\pi_\phi(y_w|s)}{\pi^{SFT}(y_w|s)} - \lambda_{KL} \log \frac{\pi_\phi(y_l|s)}{\pi^{SFT}(y_l|s)} \right) \right]. \quad (5)$$

This objective is optimized on ϕ using standard supervised learning with a static preference dataset. The gradient increases the likelihood of preferred responses y_w and decreases the likelihood of dispreferred responses y_l , relative to the reference policy.

The growing family of DPO variants is best read economically by asking what part of the preference problem changes. IPO modifies the link/objective structure to avoid some DPO overfitting behavior (Azar et al., 2024). KTO replaces pairwise comparisons with a prospect-theory-style objective over desirable and undesirable outputs (Ethayarajh et al., 2024). Nash-LHF and SPPO treat preference optimization as a game over pairwise comparisons, which is closer to social choice when preferences are cyclic or non-transitive (Munos et al., 2023; Wu et al., 2024). MaxMin-RLHF makes the aggregation rule explicitly egalitarian (Chakraborty et al., 2024). ORPO, SimPO, and robust DPO mainly change implementation, reference-model dependence, or noisy-label bias rather than the underlying welfare question (Hong et al., 2024; Meng et al., 2024; Chowdhury et al., 2024). Self-rewarding loops move some feedback generation from humans to model judges, which lowers data-collection costs but makes the evaluator itself part of the mechanism (Yuan et al., 2024).

0.3 AI Alignment as Economics

RLHF turns alignment into an economics problem: elicit preferences, aggregate them across people, design incentives for truthful feedback, and interpret the learned object. The phrase “alignment tax” appears in the assistant-alignment literature as a warning that optimizing for helpfulness, harmlessness, and honesty can trade off against other capabilities (Askell et al., 2021); Ouyang et al. (2022) provide concrete InstructGPT evidence of benchmark regressions on tasks such as SQuAD, DROP, and HellaSwag. That evidence should not be read as a law of nature. It is an empirical frontier of a particular training recipe, model class, and evaluation set.




Figure 1: **RLHF versus DPO pipelines.** Top row: the three-stage RLHF pipeline trains a reward model from human preferences, then uses PPO to fine-tune the policy with a KL penalty. Bottom row: DPO collapses the pipeline into a single supervised learning objective over preference pairs, eliminating the explicit reward model (ghosted box).

0.3.1 Social Choice and Whose Preferences

The central welfare question is whose preferences are being aligned to. If one reward model is trained on pooled labels from a non-representative group of contractors, the resulting policy optimizes an implicit aggregation rule, not society’s preferences. Social choice supplies the vocabulary for this problem: alternatives, voters, ranking rules, impossibility theorems, strategic reports, minority protection, and interpersonal aggregation (Conitzer et al., 2024; Mishra, 2023). Distributional preference learning and axiomatic analyses show that standard BTL aggregation can encode Borda-like rules or fail desirable fairness and consistency conditions (Siththaranjan et al., 2024; Ge et al., 2024). Alternatives such as personalization, MaxMin-RLHF, Nash learning from human feedback, and maximal-lottery alignment make the aggregation rule more explicit, though some of these connections remain preprint-level rather than settled practice (Park et al., 2024; Chakraborty et al., 2024; Munos et al., 2023; Maura-Rivero et al., 2025).

0.3.2 Revealed Preference Tests of LLMs

A related literature treats language models themselves as economic agents and asks whether their choices satisfy revealed-preference restrictions. Budget-allocation experiments, consumer-preference replications, moral-choice tasks, and Bayesian decision problems find that LLMs can look rational in some narrow domains and fragile in others (Chen et al., 2023; Goli and Singh, 2024; Seror, 2024; Mu et al., 2025). These tests are useful because they translate vague alignment claims into falsifiable restrictions, such as GARP consistency or proximity to a Bayes rule. But the results are prompt-, model-, version-, and domain-sensitive. They should be used as diagnostics of a trained system, not as evidence that an LLM has a stable utility function in the welfare-economics sense.

0.3.3 Mechanism Design for Feedback

Feedback is not passively observed; it is elicited through a mechanism. Labelers may be inattentive, strategic, ideologically motivated, or selected by the platform’s contracting process. Mechanism design therefore enters before the reward model is trained. VCG-style fine-tuning formulations ask how to aggregate multiple reward models when participants can benefit from manipulating the training rule (Sun et al., 2024). Strategyproof RLHF studies the same issue for preference reports, with a tradeoff between incentive compatibility and the quality of the final policy (Kleine Buening et al., 2025). Older peer-prediction mechanisms such as Bayesian truth serum and output-agreement methods show how payments can elicit truthful subjective signals when no ground truth is immediately observable (Prelec, 2004; Miller et al., 2005). For RLHF, this means the data-generating process is not a nuisance detail; it determines which preferences are observed and how costly truthful feedback is.

0.3.4 Structural Identification

Finally, preference optimization is not welfare analysis unless the learned reward has a structural interpretation. Bradley-Terry and DPO identify reward differences only through observed pairwise comparisons, up to normalizations and only on the covered comparison support. IRL identification results show that many rewards can rationalize the same behavior or preferences, and that entropy regularization, extra environments,

or downstream invariance restrictions are needed to recover sharper objects (Cao et al., 2021; Skalse et al., 2023; Knox et al., 2022). Economically, the genealogy runs from McFadden logit choice (McFadden, 1974), through Rust-style dynamic discrete choice (Rust, 1994, 1996), to maximum-entropy IRL (Ziebart et al., 2008), and finally to DPO’s log-odds reparameterization (Rafailov et al., 2023). Recent econometric work makes this bridge explicit, but still as an emerging working-paper literature rather than a settled textbook synthesis (van der Laan et al., 2025). The practical lesson is conservative: RLHF can produce useful policies, but the reward object is partially identified, aggregation-dependent, and limited by the offline support of the comparison data.

0.4 Simulation Study: Preference Learning in Job Search

A worker searches for jobs in a labor market with compensating differentials, following a McCall (1970)-style search model. Wages are measured in thousands and take values $w \in \{20, 28, 38, 50, 65, 82, 100, 125\}$, while amenities take values $z \in \{0, 1, \dots, 6\}$ and capture commute quality, flexibility, and job security. The state space has 112 states: 56 searching states in which the worker observes a pending offer (w_i, z_j) and decides to accept or reject, plus 56 employed states (w_i, z_j) in which the worker decides to stay or quit. The offer distribution exhibits compensating differentials, with wage rank and amenity rank negatively correlated ($\rho = -0.74$): high-wage offers cluster with low amenities and vice versa. The worker’s true per-period utility is $u(w, z) = \alpha \log(w) + (1 - \alpha)z$ with $\alpha = 0.6$, but this function is unobserved; the worker can only compare career trajectories (“I prefer path A to path B”), exactly as in stated-preference surveys in labor economics. While searching, the worker receives the unemployment benefit $u_b = \alpha \log(b)$ where $b = 28$. Layoffs occur with probability $p = 0.05$ per period, and the discount factor is $\gamma = 0.95$. Dynamic programming gives $V^*(s_0) = 74.13$; the optimal policy accepts 25 of 56 offer types and stays employed at 25 of 56 job types. Preference data is generated by rolling out a uniform random policy from a random searching state. Each rollout produces a career segment of $L = 15$ periods recording states and actions. For each of K comparisons, two independent career segments are generated; the segment with higher cumulative discounted utility under the true (unobserved) utility function is labeled as preferred via the Bradley-Terry model. Figure 2 shows the optimal accept/reject and stay/quit boundaries.




Figure 2: **Optimal accept/reject and stay/quit boundaries in wage-amenity space.** Left: the optimal accept/reject boundary for searching states. Right: the optimal stay/quit boundary for employed states. Both boundaries cut diagonally, reflecting the tradeoff between wage and amenity quality under compensating differentials.

Six methods are compared across $K \in \{25, 50, 100, 200, 500, 1,000, 2,000, 5,000\}$, averaged over 30 seeds. The neural network RLHF reward model is a two-layer MLP trained by Bradley-Terry MLE on the K segment pairs; per-transition rewards are discount-weighted and summed to obtain a segment score, and the resulting 112-state reward table is solved by value iteration.⁶ The correctly specified structural model parameterizes utility as $\hat{u} = \hat{\alpha} \log(w) + (1 - \hat{\alpha})z$, estimating the single parameter $\hat{\alpha}$ via Bradley-Terry MLE, then solves the MDP. The misspecified model uses $\hat{u} = \hat{\alpha} \log(w) + (1 - \hat{\alpha})\bar{z}$, where $\bar{z} = 3$ is the mean amenity level; this model ignores amenity variation across jobs, treating all amenities as identical. DPO trains a tabular softmax policy directly from trajectory comparisons, bypassing reward modeling entirely.⁷ Tabular Q-learning (10,000 episodes, $\varepsilon = 0.15$, learning rate 0.1) and exact DP provide scalar-reward baselines. All four preference methods receive identical comparison data per seed; DPO receives a same-state variant generated from the same seed.

Figure 3 reports the main results. Three findings stand out. First, the correctly specified structural model dominates: even at $K = 25$ it achieves 99.9% of DP-optimal,

⁶The neural network has 4 inputs (normalized log-wage, normalized amenity, employment indicator, action), 32 hidden units per layer, and $\sim 1,200$ parameters. The logistic loss from Equation (1) is applied to discount-weighted segment reward sums. The network parameterises $r_\theta(s, a)$, but the table handed to value iteration averages over actions to give $\bar{r}(s) = \frac{1}{2}(r_\theta(s, 0) + r_\theta(s, 1))$; this is harmless here because the true per-period reward in this environment is action-independent (actions affect only the next-state transition), so the Bradley-Terry MLE has no incentive to fit action-dependent payoffs.

⁷DPO uses 112 logit parameters ϕ_s , one per state, trained via the DPO loss (Equation 5) with Adam optimization over a sweep of $\lambda_{KL} \in \{0.01, 0.05, 0.1, 0.5, 1.0, 5.0\}$, selecting the λ_{KL} that minimizes training loss. The reference policy is uniform: $\pi^{SFT}(a|s) = 0.5$. To match the LLM setup where both completions condition on the same prompt, DPO comparison pairs start from the same initial state.




Figure 3: **Policy value $V^\pi(s_0)$ versus number of preference comparisons K for all six methods (30 seeds, $L = 15$).** The right axis shows the percentage of DP-optimal value.

illustrating the power of correct specification when the model has a single free parameter. Second, the neural network converges more slowly but reaches 99.9% by $K = 5,000$, reflecting the higher sample complexity of a flexible model relative to a one-parameter structural specification. Third, DPO plateaus at approximately 95% by $K = 500$ and does not improve further.⁸ This plateau is qualitatively different from the gridworld failure (-118%): DPO recovers a reasonable policy, but the gap does not close with additional data.⁹ The misspecified constant-amenity model plateaus at 91% regardless of K , because additional preference data cannot correct the omitted variable.

Table 1: **Diagnostics at $K = 5,000$ (single seed): policy agreement with π^* , value-function correlation, and mean accepted wage and amenity for each method.**

Method	Policy agree. (%)	V^π corr.	Mean amenity	Mean wage	$\hat{\alpha}$
NN RLHF	96.4	1.000	4.67	73	—
DPO	57.1	0.777	3.38	63	—
Correct	100.0	1.000	4.84	71	0.597
Misspecified	50.0	0.889	3.00	70	—
Optimal	100.0	1.000	4.84	71	—

Table 1 provides state-level diagnostics at $K = 5,000$. The structural model and neural

⁸DPO learns only from (s, a) pairs visited in training trajectories generated by the random behavioral policy. It cannot propagate value to undervisited states the way value iteration does after learning a reward model, so states poorly covered by the behavioral policy remain suboptimal regardless of K .

⁹DPO fails catastrophically in gridworld because transitions are stochastic (10% slip probability) and rewards are transition-dependent. The same (s, a) pair yields different rewards depending on whether the agent slipped, so the DPO loss conflates policy quality with transition luck. In the job search model, accept/reject deterministically changes employment status, and only the 5% layoff probability introduces stochastic transitions.

network achieve near-perfect policy agreement with π^* (100% and 96% respectively). DPO agrees on only 57% of states with value-function correlation 0.78, systematically underselecting on both wage and amenity dimensions.¹⁰ The misspecified model agrees on 50%, with disagreements concentrated where amenity variation matters.¹¹




Figure 4: **Policy value $V^\pi(s_0)$ versus segment length L at $K = 2,000$ for the neural network and DPO (20 seeds).** The right axis shows the percentage of DP-optimal value.

Figure 4 reports the segment length ablation at $K = 2,000$. At $L = 1$, both methods perform poorly because single-step comparisons carry minimal information about long-run value. The neural network recovers rapidly (by $L = 3$) because the reward model aggregates per-transition estimates over longer segments and value iteration propagates them through the full transition structure; DPO improves monotonically but plateaus at its $\sim 95\%$ ceiling, as it must learn the policy directly from trajectory comparisons without access to the transition model.

Two-stage RLHF separates preference estimation (a static econometric problem) from dynamic programming (which exploits the known transition model); DPO conflates the two and forfeits the transition structure that structural modelers typically have access to.¹² The DPO plateau in the job-search study is the computational version of the

¹⁰DPO’s mean accepted amenity of 3.4 parallels the misspecified structural model’s 3.0, though the mechanisms differ: the misspecified model ignores amenity variation by construction, while DPO underweights amenities because the random behavioral policy underrepresents high-amenity employed states in the training data.

¹¹An online versus offline ablation for the neural network at $K = 1,000$ (20 seeds) shows comparable performance: online 73.92 ± 0.05 , offline 73.99 ± 0.03 ($p = 0.09$). The random behavioral policy already provides diverse career trajectories covering the full wage-amenity space.

¹²This advantage is specific to settings where the transition model is known or estimable; in domains without a tractable transition model, DPO’s single-stage approach avoids compounding errors from reward model estimation.

identification point above. More comparisons tighten the empirical binary choice model on the support of observed trajectories, but they do not reveal counterfactual value in poorly covered states, nor do they identify a welfare object independent of the aggregation rule used to construct the labels.

References

- Amanda Askell, Yuntao Bai, Anna Chen, Dawn Drain, Deep Ganguli, Tom Henighan, Andy Jones, Nicholas Joseph, Ben Mann, Nova DasSarma, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Jackson Kernion, Kamal Ndousse, Catherine Olsson, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A general language assistant as a laboratory for alignment. *arXiv preprint arXiv:2112.00861*, 2021.
- Mohammad Gheshlaghi Azar, Mark Rowland, Bilal Piot, Daniel Guo, Daniele Calandriello, Michal Valko, and Remi Munos. A general theoretical paradigm to understand learning from human preferences. In *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, 2024.
- Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. The method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Haoyang Cao, Samuel N. Cohen, and Lukasz Szpruch. Identifiability in inverse reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 34, 2021.
- Souradip Chakraborty, Jiahao Qiu, Hui Yuan, Alec Koppel, Dinesh Manocha, Furong Huang, Amrit Singh Bedi, and Mengdi Wang. MaxMin-RLHF: Alignment with diverse human preferences. In *Proceedings of the 41st International Conference on Machine Learning*, 2024.
- Yiting Chen, Tracy Xiao Liu, You Shan, and Songfa Zhong. The emergence of economic rationality of GPT. *Proceedings of the National Academy of Sciences*, 120(51):e2316205120, 2023. doi: 10.1073/pnas.2316205120.
- Sayak Ray Chowdhury, Anush Kini, and Nagarajan Natarajan. Provably robust DPO: Aligning language models with noisy feedback. *arXiv preprint arXiv:2403.00409*, 2024.
- Paul F. Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

- Vincent Conitzer, Rachel Freedman, Jobst Heitzig, Wesley H. Holliday, Bob M. Jacobs, Nathan Lambert, Milan Mosse, Eric Pacuit, Stuart Russell, Hailey Schoelkopf, Emanuel Tewolde, and William S. Zwicker. Position: Social choice should guide AI alignment in dealing with diverse human feedback. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 9346–9360, 2024.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. KTO: Model alignment as prospect theoretic optimization. *arXiv preprint arXiv:2402.01306*, 2024.
- Luise Ge, Daniel Halpern, Evi Micha, Ariel D. Procaccia, Itai Shapira, Yevgeniy Vorobeychik, and Junlin Wu. Axioms for AI alignment from human feedback. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Ali Goli and Amandeep Singh. Frontiers: Can large language models capture human preferences? *Marketing Science*, 43(4):709–722, 2024. doi: 10.1287/mksc.2023.0306.
- Jiwoo Hong, Noah Lee, and James Thorne. ORPO: Monolithic preference optimization without reference model. *arXiv preprint arXiv:2403.07691*, 2024.
- Fedor Iskhakov, John Rust, and Bertel Schjerning. Machine learning and structural econometrics: Contrasts and synergies. *The Econometrics Journal*, 23(S1):S81–S124, 2020.
- Thomas Kleine Buening, Jiarui Gan, Debmalya Mandal, and Marta Kwiatkowska. Strategyproof reinforcement learning from human feedback. *arXiv preprint arXiv:2503.09561*, 2025.
- W. Bradley Knox, Stephane Hatgis-Kessel, Serena Booth, Scott Niekum, Peter Stone, and Alessandro Allievi. Models of human preference for learning reward functions. *arXiv preprint arXiv:2206.02231*, 2022.
- Tomasz Korbak, Ethan Perez, and Christopher L. Buckley. RL with KL penalties is better viewed as Bayesian inference. *arXiv preprint arXiv:2205.11275*, 2022.
- Roberto-Rafael Maura-Rivero, Marc Lanctot, Francesco Visin, and Kate Larson. Jackpot! alignment as a maximal lottery. *arXiv preprint arXiv:2501.19266*, 2025.

- Daniel McFadden. Conditional logit analysis of qualitative choice behavior. In Paul Zarembka, editor, *Frontiers in Econometrics*, pages 105–142. Academic Press, 1974. Working paper circulated 1973.
- Yu Meng, Mengzhou Xia, and Danqi Chen. SimPO: Simple preference optimization with a reference-free reward. *arXiv preprint arXiv:2405.14734*, 2024.
- Nolan Miller, Paul Resnick, and Richard Zeckhauser. Eliciting informative feedback: The peer-prediction method. *Management Science*, 51(9):1359–1373, 2005.
- Abhilash Mishra. AI alignment and social choice: Fundamental limitations and policy implications. *arXiv preprint arXiv:2310.16048*, 2023.
- Tianshi Mu, Pranjal Rawat, John Rust, Chengjun Zhang, and Qixuan Zhong. Who is more bayesian: Humans or ChatGPT? *arXiv preprint arXiv:2504.10636*, 2025.
- Remi Munos, Michal Valko, Daniele Calandriello, Mohammad Gheshlaghi Azar, Mark Rowland, Daniel Guo, Yunhao Tang, Matthieu Geist, Thomas Mesnard, Come Fiegel, Andrea Michi, Marco Selvi, Sertan Girgin, Nikola Momchev, Olivier Bachem, Daniel J. Mankowitz, Doina Precup, and Bilal Piot. Nash learning from human feedback. *arXiv preprint arXiv:2312.00886*, 2023.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744, 2022.
- Chanwoo Park, Mingyang Liu, Dingwen Kong, Kaiqing Zhang, and Asuman E. Ozdaglar. RLHF from heterogeneous feedback via personalization and preference aggregation. *arXiv preprint arXiv:2405.00254*, 2024.
- Drazen Prelec. A bayesian truth serum for subjective data. *Science*, 306(5695):462–466, 2004.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

- John Rust. Structural estimation of Markov decision processes. In Robert F. Engle and Daniel McFadden, editors, *Handbook of Econometrics*, volume 4, chapter 51, pages 3081–3143. Elsevier, 1994.
- John Rust. Numerical dynamic programming in economics. In Hans M. Amman, David A. Kendrick, and John Rust, editors, *Handbook of Computational Economics*, volume 1, pages 619–729. Elsevier, 1996.
- Avner Seror. The moral mind(s) of large language models. *arXiv preprint arXiv:2412.04476*, 2024.
- Anand Siththaranjan, Cassidy Laidlaw, and Dylan Hadfield-Menell. Distributional preference learning: Understanding and accounting for hidden context in RLHF. In *International Conference on Learning Representations*, 2024.
- Joar Max Viktor Skalse, Matthew Farrugia-Roberts, Stuart Russell, Alessandro Abate, and Adam Gleave. Invariance in policy optimisation and partial identifiability in reward learning. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 32033–32058, 2023.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F. Christiano. Learning to summarize with human feedback. In *Advances in Neural Information Processing Systems*, volume 33, pages 3008–3021, 2020.
- Haoran Sun, Yurong Chen, Siwei Wang, Wei Chen, and Xiaotie Deng. Mechanism design for LLM fine-tuning with multiple reward models. *arXiv preprint arXiv:2405.16276*, 2024.
- Lars van der Laan, Aurelien Bibaut, and Nathan Kallus. Efficient inference for inverse reinforcement learning and dynamic discrete choice models. *arXiv preprint arXiv:2512.24407*, 2025.
- Yue Wu, Zhiqing Sun, Huizhuo Yuan, Kaixuan Ji, Yiming Yang, and Quanquan Gu. Self-play preference optimization for language model alignment. *arXiv preprint arXiv:2405.00675*, 2024.

Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. Self-rewarding language models. *arXiv preprint arXiv:2401.10020*, 2024.

Brian D. Ziebart, Andrew Maas, J. Andrew Bagnell, and Anind K. Dey. Maximum entropy inverse reinforcement learning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2008.

Daniel M. Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B. Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.